

Programming Machine Learning Applications

Lecture Six: Regression and Gradient Descent

Dr. Aleksandar Velkoski



Text Classification

Recommender Systems

Review of Lecture Five

Regression

Gradient Descent

Discuss Code

Lecture Six

Regression

What is Numeric Prediction

(Numerical) prediction is similar to classification

- construct a model
- use model to predict value for a given input

Prediction is different from classification

- Classification refers to predicting categorical class label
- Prediction models continuous-valued functions

Major method for prediction: regression

- model the relationship between one or more independent or predictor variables and a dependent or response variable


Regression analysis

- Linear and multiple regression
- Non-linear regression
- Other regression methods: generalized linear model, Poisson regression, log-linear models, regression trees

Simple Linear Regression


Simple linear regression: involves a response variable y and a single predictor variable x where $y = w_0 + w_1x$

| x | y |
|----------------------------------|-----------------|
| Living area (feet ²) | Price (1000\$s) |
| 2104 | 400 |
| 1600 | 330 |
| 2400 | 369 |
| 1416 | 232 |
| 3000 | 540 |
| \vdots | \vdots |



Goal: Use data to estimate weights (parameters) w_0 and w_1 such that prediction error is minimized

Simple Linear Regression



Simple Linear Regression

i Method of least squares:

- 4 Estimates the best-fitting straight line
- 4 w_0 and w_1 are obtained by minimizing the sum of the squared errors (a.k.a. residuals)

$$\begin{aligned}SSE &= \sum_i e_i^2 = \sum_i (y_i - \hat{y}_i)^2 \\&= \sum_i (y_i - (w_0 + w_1 x_i))^2\end{aligned}$$

Multiple Linear Regression

i **Multiple linear regression**: involves more than one predictor variable

- 4 Features represented as x_1, x_2, \dots, x_d
- 4 Training data is of the form $(\mathbf{x}^1, y^1), (\mathbf{x}^2, y^2), \dots, (\mathbf{x}^n, y^n)$
(each \mathbf{x}^j is a row vector in matrix \mathbf{X} , i.e. a row in the data)
- 4 For a specific value of a feature x_i in data item \mathbf{x}^j we use: x_i^j
- 4 Ex. For 2-D data, the regression function is: $\hat{y} = w_0 + w_1x_1 + w_2x_2$

| x_1 | x_2 | y |
|----------------------------------|-----------|-----------------|
| Living area (feet ²) | #bedrooms | Price (1000\$s) |
| 2104 | 3 | 400 |
| 1600 | 3 | 330 |
| 2400 | 3 | 369 |
| 1416 | 2 | 232 |
| 3000 | 4 | 540 |
| \vdots | \vdots | \vdots |

Least Squares Generalization

i Multiple dimensions

4 To simplify add a new feature $x_0 = \mathbf{1}$ to feature vector \mathbf{x} :

| \mathbf{x}_0 | \mathbf{x}_1 | \mathbf{x}_2 | \mathbf{y} |
|----------------|----------------------------------|----------------|----------------|
| | Living area (feet ²) | #bedrooms | Price (1000\$) |
| 1 | 2104 | 3 | 400 |
| 1 | 1600 | 3 | 330 |
| 1 | 2400 | 3 | 369 |
| 1 | 1416 | 2 | 232 |
| 1 | 3000 | 4 | 540 |
| \vdots | \vdots | \vdots | \vdots |

$$\hat{y} = f(x_0, x_1, \dots, x_d) = w_0 x_0 + \sum_{i=1}^d w_i x_i = \sum_{i=0}^d w_i x_i = \mathbf{w}^T \cdot \mathbf{x}$$

Least Squares Generalization

$$\hat{y} = f(x_0, x_1, \dots, x_d) = f(\mathbf{x}) = w_0 x_0 + \sum_{i=1}^d w_i x_i = \sum_{i=0}^d w_i x_i = \mathbf{w}^T \cdot \mathbf{x}$$


4 Calculate the error function (SSE) and determine \mathbf{w} :

$$E(\mathbf{w}) = (\mathbf{y} - f(\mathbf{x}))^2 = \left\| \mathbf{y} - \sum_{i=0}^d w_i \cdot \mathbf{x}_i \right\|^2 = \sum_{j=1}^n (y^j - \sum_{i=0}^d w_i \cdot x_i^j)^2$$

Application of Linear Regression

- i **The inputs X for linear regression can be:**
 - 4 Original quantitative inputs
 - 4 Transformation of quantitative inputs, e.g. log, exp, square root, square, etc.
 - 4 Polynomial transformation
 - h example: $y = w_0 + w_1 \cdot x + w_2 \cdot x^2 + w_3 \cdot x^3$
 - 4 Dummy coding of categorical inputs
 - 4 Interactions between variables
 - h example: $x_3 = x_1 \cdot x_2$
- i **This allows use of linear regression techniques to fit much more complicated non-linear datasets.**

Fitting Polynomial with Linear Model



$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

Regularization

- i Complex models (lots of parameters) are often prone to overfitting
- i Overfitting can be reduced by imposing a constraint on the overall magnitude of the parameters (i.e., by including coefficients as part of the optimization process)
- i Two common types of regularization in linear regression:
 - 4 L_2 regularization (a.k.a. **ridge regression**). Find \mathbf{w} which minimizes:

$$\sum_{j=1}^N (y_j - \sum_{i=0}^d w_i \cdot x_i)^2 + \lambda \sum_{i=1}^d w_i^2$$

h λ is the regularization parameter: bigger λ imposes more constraint

- 4 L_1 regularization (a.k.a. **lasso**). Find \mathbf{w} which minimizes:

$$\sum_{j=1}^N (y_j - \sum_{i=0}^d w_i \cdot x_i)^2 + \lambda \sum_{i=1}^d |w_i|$$

Other Regression-Based Models

Generalized linear models

- Foundation on which linear regression can be applied to modeling categorical response variables
- Variance of y is a function of the mean value of y , not a constant
- Logistic regression: models the probability of some event occurring as a linear function of a set of predictor variables
- Poisson regression: models the data that exhibit a Poisson distribution

Log-linear models (for categorical data)

- Approximate discrete multidimensional prob. distributions
- Also useful for data compression and smoothing

Regression trees and model trees

- Trees to predict continuous values rather than class labels

Regression Trees & Model Trees

Regression tree: proposed in CART system (Breiman et al. 1984)

- CART: Classification And Regression Trees
- Each leaf stores a continuous-valued prediction
- It is the average value of the predicted attribute for the training instances that reach the leaf

Model tree: proposed by Quinlan (1992)

- Each leaf holds a regression model—a multivariate linear equation for the predicted attribute
- A more general case than regression tree

Regression and model trees tend to be more accurate than linear regression when instances are not represented well by simple linear models

Evaluating Regression Models

i **Prediction Accuracy**

- 4 Difference between predicted scores and the actual results (from evaluation set)
- 4 Typically the accuracy of the model is measured in terms of variance (i.e., average of the squared differences)

i **Common Metrics** (p_i = predicted target value for test instance i , a_i = actual target value for instance i)


- 4 **Mean Absolute Error:** Average loss over the test set

$$MAE = \frac{(p_1 - a_1) + \dots + (p_n - a_n)}{n}$$


- 4 **Root Mean Squared Error:** compute the standard deviation (i.e., square root of the co-variance between predicted and actual ratings)

$$RMSE = \sqrt{\frac{(p_1 - a_1)^2 + \dots + (p_n - a_n)^2}{n}}$$


Example: Web Traffic




Example: Web Traffic 1d Fit




Example: Web Traffic 2d Fit




Example: Web Traffic Beyond 2d Fit



Model Complexity



Bias-Variance Tradeoff



Bias-Variance Tradeoff

i Possible ways of dealing with high bias

- 4 Get additional features
- 4 More complex model (e.g., adding polynomial terms such as x_1^2 , x_2^2 , x_1x_2 , etc.)
- 4 Use smaller regularization coefficient λ .
- 4 **Note:** getting more training data won't necessarily help in this case

i Possible ways dealing with high variance

- 4 Use more training instances
- 4 Reduce the number of features
- 4 Use simpler models
- 4 Use a larger regularization coefficient λ .

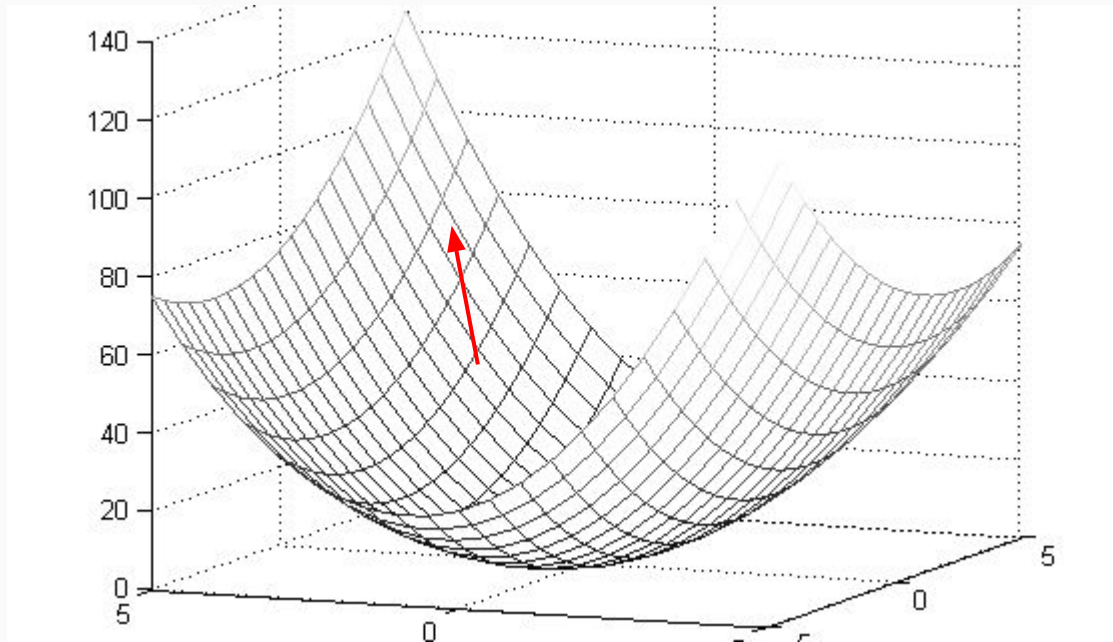
Gradient Descent

Gradient Descent Optimization

- i **Linear regression can also be solved using Gradient Decent optimization approach**
- i **GD can be used in a variety of settings to find the minimum value of functions (including non-linear functions) where a closed form solution is not available or not easily obtained**
- i **Basic idea:**
 - 4 Given an objective function $J(\mathbf{w})$ (e.g., sum of squared errors), with \mathbf{w} as a vector of variables w_0, w_1, \dots, w_d , iteratively minimize $J(\mathbf{w})$ by **finding the gradient** of the function surface in the variable-space and **adjusting the weights** in the opposite direction
 - 4 The gradient is a vector with each element representing the slope of the function in the direction of one of the variables
 - 4 Each element is the partial derivative of function with respect to one of variables

$$\nabla J(\mathbf{w}) = \nabla J(w_1, w_2, \dots, w_d) = \begin{bmatrix} \frac{\partial f(\mathbf{w})}{\partial w_1} & \frac{\partial f(\mathbf{w})}{\partial w_2} & \dots & \frac{\partial f(\mathbf{w})}{\partial w_d} \end{bmatrix}$$

Gradient Descent Optimization




Gradient is a Vector

$$\frac{\partial}{\partial m} = \frac{2}{N} \sum_{i=1}^N -x_i(y_i - (mx_i + b))$$

$$\frac{\partial}{\partial b} = \frac{2}{N} \sum_{i=1}^N -(y_i - (mx_i + b))$$

Gradient Descent Optimization

Gradient vector points in the direction of the steepest ascent of function



Gradient Descent Optimization

- i **Finding minimum directly by closed form analytical solution often difficult or impossible**

- 4 Quadratic functions in many variables

- h system of equations for partial derivatives may be ill-conditioned

- h example: linear least squares fit where redundancy among features is high

- 4 Other convex functions

- h global minimum exists, but there is no closed form solution

- h example: maximum likelihood solution for logistic regression

- 4 Nonlinear functions

- h partial derivatives are not linear

- h example: $f(x_1, x_2) = x_1(\sin(x_1x_2)) + x_2^2$

- h example: sum of transfer functions in neural networks

Gradient Descent Optimization


i **Given an objective (e.g., error) function** $E(\mathbf{w}) = E(w_0, w_1, \dots, w_d)$

i **Process (follow the gradient *downhill*):**

1. Pick an initial set of weights (random): $\mathbf{w} = (w_0, w_1, \dots, w_d)$
2. Determine the descent direction: $-\nabla E(\mathbf{w}^t)$
3. Choose a learning rate: η
4. Update your position: $\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \cdot \nabla E(\mathbf{w}^t)$
5. Repeat from 2) until stopping criterion is satisfied

i **Typical stopping criteria**

- 4 $\nabla E(\mathbf{w}^{t+1}) \sim 0$
- 4 some validation metric is optimized



Note: this step involves simultaneous updating of each weight w_i

Gradient Descent Optimization

i **In Least Squares Regression:** $E(\mathbf{w}) = \left\| \mathbf{y} - \sum_{i=0}^d w_i \cdot \mathbf{x}_i \right\|^2 = (\mathbf{y} - \mathbf{w}^\top \cdot \mathbf{x})^2$

i **Process (follow the gradient *downhill*):**


1. Select initial $\mathbf{w} = (w_0, w_1, \dots, w_d)$

2. Compute $-\nabla E(\mathbf{w})$

3. Set η


4. Update: $\mathbf{w} := \mathbf{w} - \eta \cdot \nabla E(\mathbf{w})$

5. Repeat until $\nabla E(\mathbf{w}^{t+1}) \sim 0$



$$w_j := w_j - \eta \frac{1}{2n} \sum_{i=1}^n (\mathbf{w}^\top \cdot \mathbf{x}^i - y^i) x_j^i$$

for $j = 0, 1, \dots, d$


Illustration




Illustration



Illustration



Illustration



Gradient Descent Optimization

i Problems:


4 Choosing step size (learning rate)

h too small \rightarrow convergence is slow and inefficient

h too large \rightarrow may not converge

4 Can get stuck on “flat” areas of function

4 Easily trapped in local minima



Stochastic Gradient Descent

- i **Application to training a machine learning model:**

1. Choose one sample from training set: x^i
2. Calculate objective function for that single sample: $(\mathbf{w}^T \cdot \mathbf{x}^i - y^i)^2$
3. Calculate gradient from objective function:
4. Update model parameters a single step based on gradient and learning rate:

$$w_j := w_j - \eta(\mathbf{w}^T \cdot \mathbf{x}^i - y^i)x_j^i \quad \text{for } j = 0, \dots, d$$

5. Repeat from 1) until stopping criterion is satisfied

- i **Typically entire training set is processed multiple times before stopping**
- i **Order in which samples are processed can be fixed or random.**

Code

Wrapping-up the Lecture

Questions

What is the goal of Linear Regression?

Describe two types of Regularization. What's the intuition of each type?

Why is getting trapped in a local minima a problem in Gradient Descent?